


# CSCI4333 Database Design & Implement

## **Lecture Ten – Relational Model**

### **4**

Instructor: Dr. Yifeng Gao

# From Relational Schema to Table



# Foreign Keys, Referential Integrity

- Foreign key : Set of fields in one relation that is used to 'refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a 'logical pointer'.
- e.g. sid is a foreign key referring to **Students**:
  - Enrolled(sid: string, *crn*: string)
  - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.

# Foreign Keys

- Only students listed in the Students relation should be allowed to enroll for courses.

## Enrolled

sid	cid	grade
<u>53666</u>	Carnatic101	C
<u>53666</u>	Reggae203	B
53650	Topology112	A
<u>53666</u>	History105	B


## Students

sid	name	login	age	gpa
<u>53666</u>	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
<u>53650</u>	Smith	smith@math	19	3.8

# Quick Question


- With foreign key constraint, can I add
  - (51111, History105,A) to Enroll Table? ✗
  - (51111, John, john@cs, 17, 3.5) to Student Table? ✓

Enrolled



sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students



sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

# Enforcing Referential Integrity

- Consider Students and Enrolled; sid in Enrolled is a foreign key that references Students.
- What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a Students tuple is deleted?
  - Also delete all Enrolled tuples that refer to it.
  - Disallow deletion of a Students tuple that is referred to.
  - Set sid in Enrolled tuples that refer to it to a *default sid*.
  - (In SQL, also: Set sid in Enrolled tuples that refer to it to a special value null, denoting 'unknown' or 'inapplicable'.)

# Creating Relations in SQL

Students(sid: string, name: string)

- Creates a Students relation.
  - Observe that the type (domain) of each field is specified
  - enforced by the DBMS whenever tuples are added or modified.

CREATE TABLE Students(  
    sid CHAR(20),  
    name CHAR(20),  
    PRIMARY KEY sid  
);

# Creating Relations in SQL

Course(Name: string, CRN: string)

```
CREATE TABLE Course(  
  CRN CHAR(20),  
  name CHAR(20),  
  PRIMARY KEY CRN  
);
```

Enroll(sid: string, CRN: string)

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
  crn CHAR(20),  
  PRIMARY KEY (sid,cid),  
  FOREIGN KEY (sid) REFERENCES Students,  
  FOREIGN KEY (crn) REFERENCES Course  
);
```



# Primary and Candidate Keys in SQL

- Possibly many candidate keys (specified using UNIQUE), one of which is chosen as the *primary key*.
- Suppose a table only have one candidate key:

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid,cid) );
```

# Primary and Candidate Keys in SQL


- Possibly many candidate keys (specified using **UNIQUE**), one of which is chosen as the *primary key*.
- Suppose a table only have **two candidate keys**:

```
CREATE TABLE Car
  (VIN CHAR(20),
   License# CHAR(20),
   State CHAR(2),
   PRIMARY KEY (VIN),
   UNIQUE (License#, State)
);
```

# Translating ER Diagrams with Key Constraints

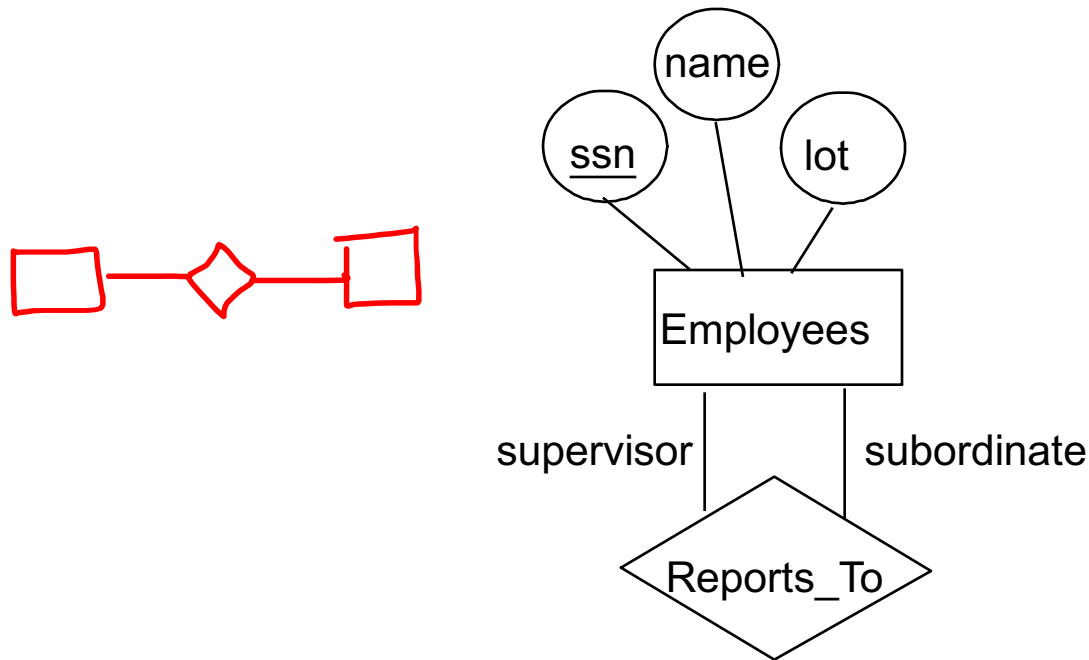
- Option 1: Map relationship to a table:
  - Note that *did* is the key now!
  - Separate tables for Employees and Departments.
- Option 2: Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Manager,  
  FOREIGN KEY (did) REFERENCES Departments);
```



```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Manager);
```


# Relationship with Roles



```
CREATE TABLE Reports_To(  
  supervisor_ssn CHAR(11),  
  subordinate_ssn CHAR(11),  
  PRIMARY KEY (supervisor_ssn, subordinate_ssn),  
  FOREIGN KEY (supervisor_ssn) REFERENCES Employees(ssn),  
  FOREIGN KEY (subordinate_ssn) REFERENCES Employees(ssn));
```

# Participation Constraints

- Does every department have a manager?
  - If so, this is a *participation constraint*: the participation of Departments in Manages is said to be *total* (vs. *partial*).
  - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)




# Participation Constraints in SQL

- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees);
```

# Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
  - Weak entity set must have total participation in this *identifying* relationship set.



# Translating Weak Entity Sets

- Weak entity set and identifying relationship set are translated into a single table.
  - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```



# Referential Integrity in SQL

- SQL/92 and SQL:1999 support all 4 options on deletes and updates.
  - Default is NO ACTION (*delete/update is rejected*)
  - CASCADE (also delete all tuples that refer to deleted tuple)
  - SET NULL / SET DEFAULT (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

# Relational Model: Summary

- A tabular representation of data.
- Simple and intuitive, currently the most widely used.
- Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
  - Two important ICs: primary and foreign keys
  - In addition, we *always* have domain constraints.
- Powerful and natural query languages exist.
- Rules to translate ER to relational model